

# 01 — Exploratory Data Analysis (EDA)

**Dataset:** Kaggle Credit Card Fraud Detection (`creditcard.csv`)

**Goal:** Understand structure, class imbalance, and basic patterns that differentiate *fraud* vs *legit* transactions.

Key checks:

1. Schema overview (rows, columns, dtypes, missingness)
2. Class imbalance (0=legit, 1=fraud)
3. Descriptives by class
4. Distributions: `Amount`, `Time`, selected `V*` components
5. Simple correlations with `Class`
6. Quick anomalies sanity checks (duplicates, extreme values)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

pd.set_option("display.precision", 3)
pd.set_option("display.max_columns", 50)

sns.set(context="notebook", style="whitegrid")
```

## Load dataset

The dataset contains 284,807 transactions and 31 columns:

`Time`, `Amount`, `Class`, and 28 anonymised PCA-like features `V1..V28`.

`Class=1` denotes *fraud*.

```
PATH = "../data/creditcard.csv"
df = pd.read_csv(PATH)
df.shape, df.columns.tolist()[:10]

((284807, 31), ['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7',
'V8', 'V9'])
```

## Schema & missingness

We inspect dtypes and nulls. The Kaggle version typically has **no missing values**, but we verify.

```
df.info()
df.isna().sum().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
--- 
 0   Time    284807 non-null   float64
 1   V1      284807 non-null   float64
 2   V2      284807 non-null   float64
 3   V3      284807 non-null   float64
 4   V4      284807 non-null   float64
 5   V5      284807 non-null   float64
 6   V6      284807 non-null   float64
 7   V7      284807 non-null   float64
 8   V8      284807 non-null   float64
 9   V9      284807 non-null   float64
 10  V10     284807 non-null   float64
 11  V11     284807 non-null   float64
 12  V12     284807 non-null   float64
 13  V13     284807 non-null   float64
 14  V14     284807 non-null   float64
 15  V15     284807 non-null   float64
 16  V16     284807 non-null   float64
 17  V17     284807 non-null   float64
 18  V18     284807 non-null   float64
 19  V19     284807 non-null   float64
 20  V20     284807 non-null   float64
 21  V21     284807 non-null   float64
 22  V22     284807 non-null   float64
 23  V23     284807 non-null   float64
 24  V24     284807 non-null   float64
 25  V25     284807 non-null   float64
 26  V26     284807 non-null   float64
 27  V27     284807 non-null   float64
 28  V28     284807 non-null   float64
 29  Amount   284807 non-null   float64
 30  Class    284807 non-null   int64  
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

0

## Class balance

The dataset is **highly imbalanced**. We compute counts and proportions.


```
counts = df['Class'].value_counts().rename({0:"legit", 1:"fraud"})
props  = (counts / counts.sum()).round(4)
display(pd.DataFrame({"count": counts, "prop": props}))
```

```

plt.figure(figsize=(4,3))
sns.barplot(x=counts.index, y=counts.values)
plt.title("Class distribution (0=legit, 1=fraud) ")
plt.ylabel("Count"); plt.xlabel("Class")
plt.tight_layout(); plt.show()

```

Class	count	prop
legit	284315	0.998
fraud	492	0.002



## Descriptive statistics (overall & by class)

We look at summary stats overall and then compare *fraud vs legit* for *Amount* and *Time*.

```

display(df.describe().T[['mean', 'std', 'min', 'max']].head(12))

byc = df.groupby('Class')
[['Amount', 'Time']].agg(['count', 'mean', 'median', 'std', 'min', 'max'])
byc

```

	mean	std	min	max
Time	9.481e+04	47488.146	0.000	172792.000
V1	1.168e-15	1.959	-56.408	2.455
V2	3.417e-16	1.651	-72.716	22.058
V3	-1.380e-15	1.516	-48.326	9.383
V4	2.074e-15	1.416	-5.683	16.875
V5	9.604e-16	1.380	-113.743	34.802
V6	1.487e-15	1.332	-26.161	73.302
V7	-5.556e-16	1.237	-43.557	120.589
V8	1.213e-16	1.194	-73.217	20.007
V9	-2.406e-15	1.099	-13.434	15.595

```

V10    2.239e-15      1.089   -24.588      23.745
V11    1.673e-15      1.021   -4.797      12.019

          Amount                                Time
\       count     mean  median      std   min      max   count
mean
Class

0      284315   88.291   22.00  250.105   0.0  25691.16  284315
94838.202
1      492     122.211   9.25  256.683   0.0  2125.87     492
80746.807

          median      std   min      max
Class
0      84711.0  47484.016   0.0  172792.0
1      75568.5  47835.365  406.0  170348.0

```

## Basic hygiene checks

- Duplicates (row-level)
- Constant/near-zero-variance columns
- Obvious invalid values (negative amounts)

```

dup_rows = df.duplicated().sum()
neg_amount = (df['Amount'] < 0).sum()
nzv = df.drop(columns=['Class']).nunique().sort_values().head(5)

print(f"Duplicate rows: {dup_rows}")
print(f"Negative Amount values: {neg_amount}")
nzv

```

## Distributions: Amount & Time

`Time` is seconds from the first transaction; `Amount` is the transaction value.  
We compare distributions for fraud vs legit using random subsamples for readability.

```

from scipy import stats

sample = df.sample(50000, random_state=425)

fig, axes = plt.subplots(1, 2, figsize=(10, 4))

# Separate the data
legit_amount = sample[sample['Class'] == 0]['Amount']
fraud_amount = sample[sample['Class'] == 1]['Amount']
legit_time = sample[sample['Class'] == 0]['Time']

```

```

fraud_time = sample[sample['Class'] == 1]['Time']

# Plot Amount KDE
x_range_amount = np.linspace(min(sample['Amount']),
max(sample['Amount']), 100)
kde_legit_amount = stats.gaussian_kde(legit_amount)
kde_fraud_amount = stats.gaussian_kde(fraud_amount)

axes[0].plot(x_range_amount, kde_legit_amount(x_range_amount),
label='Legit', alpha=0.7)
axes[0].plot(x_range_amount, kde_fraud_amount(x_range_amount),
label='Fraud', alpha=0.7)
axes[0].fill_between(x_range_amount, kde_legit_amount(x_range_amount),
alpha=0.3)
axes[0].fill_between(x_range_amount, kde_fraud_amount(x_range_amount),
alpha=0.3)
axes[0].set_title("Amount density by Class")
axes[0].legend()

# Plot Time KDE
x_range_time = np.linspace(min(sample['Time']), max(sample['Time']),
100)
kde_legit_time = stats.gaussian_kde(legit_time)
kde_fraud_time = stats.gaussian_kde(fraud_time)

axes[1].plot(x_range_time, kde_legit_time(x_range_time),
label='Legit', alpha=0.7)
axes[1].plot(x_range_time, kde_fraud_time(x_range_time),
label='Fraud', alpha=0.7)
axes[1].fill_between(x_range_time, kde_legit_time(x_range_time),
alpha=0.3)
axes[1].fill_between(x_range_time, kde_fraud_time(x_range_time),
alpha=0.3)
axes[1].set_title("Time density by Class")
axes[1].legend()

plt.tight_layout()
plt.show()

```

## Boxplots: Amount and Time by Class

Boxplots highlight differences in central tendency and spread between classes.

```

fig, axes = plt.subplots(1,2, figsize=(10,4))
sns.boxplot(data=sample, x="Class", y="Amount", ax=axes[0])
axes[0].set_title("Amount by Class")
sns.boxplot(data=sample, x="Class", y="Time", ax=axes[1])
axes[1].set_title("Time by Class")
plt.tight_layout(); plt.show()

```

## Distributions of selected v\* components

The V1..V28 columns are anonymised components (PCA-like).

We plot a few that often show separation (your results may vary): V3, V4, V10, V14.

```
cols = ["V3", "V4", "V10", "V14"]
fig, axes = plt.subplots(2, 2, figsize=(10, 7))
for col, ax in zip(cols, axes.ravel()):
    # Plot density for each class separately
    sample[sample['Class'] == 0][col].plot.density(ax=ax,
label='Legit', alpha=0.7)
    sample[sample['Class'] == 1][col].plot.density(ax=ax,
label='Fraud', alpha=0.7)
    ax.fill_between(ax.lines[0].get_xdata(), ax.lines[0].get_ydata(),
alpha=0.3)
    ax.fill_between(ax.lines[1].get_xdata(), ax.lines[1].get_ydata(),
alpha=0.3)
    ax.set_title(f"{col} density by Class")
    ax.legend()
plt.tight_layout(); plt.show()
```

## Correlation with class

Because Class is binary, Pearson correlation is equivalent to point-biserial correlation.  
We compute correlations between each feature and Class and show the top signals.

```
corr = df.corr(numeric_only=True)
['Class'].drop('Class').sort_values(key=np.abs, ascending=False)
corr.head(12).to_frame("corr_with_Class")
```

## Fraud vs Legit: side-by-side summary table

Simple per-class summaries to identify directional differences.

```
summary_cols = ["Amount", "Time", "V3", "V4", "V10", "V14"]
summary = (df.groupby("Class")[summary_cols]
            .agg(['mean', 'median', 'std', 'min', 'max'])
            .round(3))
summary
```